

Time Series Analysis & Forecasting Using R

10. Forecast reconciliation



Outline

- 1 Hierarchical and grouped time series
- 2 Forecast reconciliation
- 3 Example: Australian tourism
- 4 Lab Session 20

Outline

1 Hierarchical and grouped time series

2 Forecast reconciliation

3 Example: Australian tourism

4 Lab Session 20

Australian Pharmaceutical Benefits Scheme



PBS sales

PBS

```
# A tibble: 67,596 x 9 [1M]
```

```
# Key:      Concession, Type, ATC1, ATC2 [336]
```

	Month	Concession	Type	ATC1	ATC1_desc	ATC2	ATC2_desc	Scripts	
	<mth>	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>	<dbl>	
1	1991	Jul	Concessional	Co-pa~	A	Alimenta~	A01	STOMATOL~	18228
2	1991	Aug	Concessional	Co-pa~	A	Alimenta~	A01	STOMATOL~	15327
3	1991	Sep	Concessional	Co-pa~	A	Alimenta~	A01	STOMATOL~	14775
4	1991	Oct	Concessional	Co-pa~	A	Alimenta~	A01	STOMATOL~	15380
5	1991	Nov	Concessional	Co-pa~	A	Alimenta~	A01	STOMATOL~	14371
6	1991	Dec	Concessional	Co-pa~	A	Alimenta~	A01	STOMATOL~	15028
7	1992	Jan	Concessional	Co-pa~	A	Alimenta~	A01	STOMATOL~	11040
8	1992	Feb	Concessional	Co-pa~	A	Alimenta~	A01	STOMATOL~	15165
9	1992	Mar	Concessional	Co-pa~	A	Alimenta~	A01	STOMATOL~	16898
10	1992	Apr	Concessional	Co-pa~	A	Alimenta~	A01	STOMATOL~	18141

```
# i 67,586 more rows
```

```
# i 1 more variable: Cost <dbl>
```

ATC drug classification

- A Alimentary tract and metabolism
- B Blood and blood forming organs
- C Cardiovascular system
- D Dermatologicals
- G Genito-urinary system and sex hormones
- H Systemic hormonal preparations, excluding sex hormones and insulins
- J Anti-infectives for systemic use
- L Antineoplastic and immunomodulating agents
- M Musculo-skeletal system
- N Nervous system
- P Antiparasitic products, insecticides and repellents
- R Respiratory system
- S Sensory organs
- V Various

ATC drug classification

ATC1: 14 classes

A

Alimentary tract and metabolism

ATC2: 84 classes

A10

Drugs used in diabetes

A10B

Blood glucose lowering drugs


A10BA

Biguanides

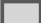



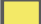



A10BA02

Metformin

Australian tourism



State

	Australian Capital Territory
	New South Wales
	Northern Territory
	Queensland
	South Australia
	Tasmania
	Victoria
	Western Australia

Australian tourism

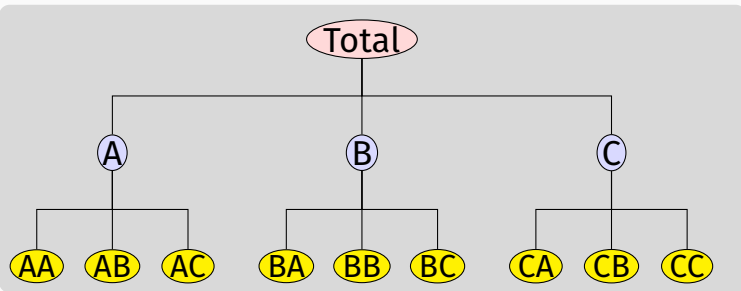
tourism

```
# A tsibble: 24,320 x 5 [1Q]
# Key:           Region, State, Purpose [304]
  Quarter Region  State Purpose  Trips
   <qtr> <chr>    <chr> <chr>    <dbl>
1 1998 Q1 Adelaide SA      Business 135.
2 1998 Q2 Adelaide SA      Business 110.
3 1998 Q3 Adelaide SA      Business 166.
4 1998 Q4 Adelaide SA      Business 127.
5 1999 Q1 Adelaide SA      Business 137.
6 1999 Q2 Adelaide SA      Business 200.
7 1999 Q3 Adelaide SA      Business 169.
8 1999 Q4 Adelaide SA      Business 134.
9 2000 Q1 Adelaide SA      Business 154.
10 2000 Q2 Adelaide SA      Business 169.
# i 24,310 more rows
```

- Quarterly data on visitor nights, 1998:Q1 – 2017:Q4
- From: *National Visitor Survey*, based on annual interviews of 120,000 Australians aged 15+, collected by Tourism Research Australia.
- Split by 8 states and 76 regions
- Split by purpose of travel
 - ▶ Holiday
 - ▶ Visiting friends and relatives (VFR)
 - ▶ Business
 - ▶ Other
- 304 bottom-level series


Hierarchical time series

A **hierarchical time series** is a collection of several time series that are linked together in a hierarchical structure.



Hierarchical time series

A **hierarchical time series** is a collection of several time series that are linked together in a hierarchical structure.

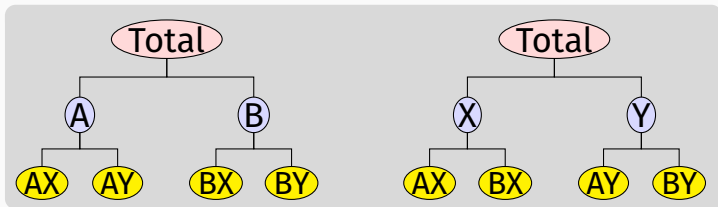


Examples

- PBS sales by ATC groups
- Tourism demand by states, regions

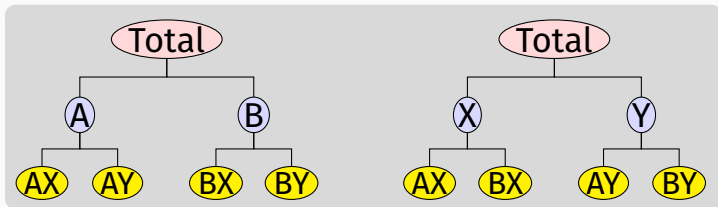
Grouped time series

A **grouped time series** is a collection of time series that can be grouped together in a number of non-hierarchical ways.



Grouped time series

A **grouped time series** is a collection of time series that can be grouped together in a number of non-hierarchical ways.



Examples

- Tourism by state and purpose of travel
- Retail sales by product groups/sub groups, and by countries/regions

Creating aggregates

```
PBS |>  
  aggregate_key(ATC1 / ATC2, Scripts = sum(Scripts)) |>  
  filter(Month == yearmonth("1991 Jul")) |>  
  print(n = 18)
```

```
# A tsibble: 98 x 4 [1M]  
# Key:      ATC1, ATC2 [98]  
  Month ATC1      ATC2      Scripts  
  <mt> <chr*>    <chr*>    <dbl>  
1 1991 Jul <aggregated> <aggregated> 8090395  
2 1991 Jul A      <aggregated> 799025  
3 1991 Jul B      <aggregated> 109227  
4 1991 Jul C      <aggregated> 1794995  
5 1991 Jul D      <aggregated> 299779  
6 1991 Jul G      <aggregated> 300931  
7 1991 Jul H      <aggregated> 112114  
8 1991 Jul J      <aggregated> 1151681  
9 1991 Jul L      <aggregated> 24580  
10 1991 Jul M      <aggregated> 562956  
11 1991 Jul N      <aggregated> 1546023  
12 1991 Jul P      <aggregated> 47661  
13 1991 Jul R      <aggregated> 859273  
14 1991 Jul S      <aggregated> 391639  
15 1991 Jul V      <aggregated> 38705  
16 1991 Jul Z      <aggregated> 51806
```

Creating aggregates

```
tourism |>
  aggregate_key(Purpose * (State / Region), Trips = sum(Trips)) |>
  filter(Quarter == yearquarter("1998 Q1")) |>
  print(n = 15)
```

```
# A tsibble: 425 x 5 [1Q]
```

```
# Key:      Purpose, State, Region [425]
```

	Quarter	Purpose	State	Region	Trips
	<qtr>	<chr*>	<chr*>	<chr*>	<dbl>
1	1998 Q1	<aggregated>	<aggregated>	<aggregated>	23182.
2	1998 Q1	Business	<aggregated>	<aggregated>	3599.
3	1998 Q1	Holiday	<aggregated>	<aggregated>	11806.
4	1998 Q1	Other	<aggregated>	<aggregated>	680.
5	1998 Q1	Visiting	<aggregated>	<aggregated>	7098.
6	1998 Q1	<aggregated>	ACT	<aggregated>	551.
7	1998 Q1	<aggregated>	NSW	<aggregated>	8040.
8	1998 Q1	<aggregated>	NT	<aggregated>	181.
9	1998 Q1	<aggregated>	QLD	<aggregated>	4041.
10	1998 Q1	<aggregated>	SA	<aggregated>	1735.
11	1998 Q1	<aggregated>	TAS	<aggregated>	982.
12	1998 Q1	<aggregated>	VIC	<aggregated>	6010.
13	1998 Q1	<aggregated>	WA	<aggregated>	1641.
14	1998 Q1	<aggregated>	ACT	Canberra	551.

Creating aggregates

- Similar to `summarise()` but using the key structure
- A grouped structure is specified using `grp1 * grp2`
- A nested structure is specified via `parent / child`.
- Groups and nesting can be mixed:

```
(country/region/city) * (brand/product)
```
- All possible aggregates are produced.
- These are useful when forecasting at different levels of aggregation.

Outline

1 Hierarchical and grouped time series

2 Forecast reconciliation

3 Example: Australian tourism

4 Lab Session 20

The problem

- 1 How to forecast time series at all nodes such that the forecasts add up in the same way as the original data?
- 2 Can we exploit relationships between the series to improve the forecasts?

The problem

- 1 How to forecast time series at all nodes such that the forecasts add up in the same way as the original data?
- 2 Can we exploit relationships between the series to improve the forecasts?

The solution

- 1 Forecast all series at all levels of aggregation using an automatic forecasting algorithm.
(e.g., ETS, ARIMA, ...)
- 2 Reconcile the resulting forecasts so they add up correctly using least squares optimization (i.e., find closest reconciled forecasts to the original forecasts).
- 3 This is available using `reconcile()`.

Forecast reconciliation

```
tourism |>
  aggregate_key(Purpose * (State / Region), Trips = sum(Trips)) |>
  model(ets = ETS(Trips)) |>
  reconcile(ets_adjusted = min_trace(ets)) |>
  forecast(h = 2)
```

```
# A fable: 1,700 x 7 [1Q]
```

```
# Key:      Purpose, State, Region, .model [850]
```

	Purpose	State	Region	.model	Quarter	Trips	.mean
	<chr*>	<chr*>	<chr*>	<chr>	<qtr>	<dist>	<dbl>
1	Business	ACT	Canberra	ets	2018 Q1	N(144, 1119)	144.
2	Business	ACT	Canberra	ets	2018 Q2	N(203, 2260)	203.
3	Business	ACT	Canberra	ets_adjusted	2018 Q1	N(157, 539)	157.
4	Business	ACT	Canberra	ets_adjusted	2018 Q2	N(214, 951)	214.
5	Business	ACT	<aggregated>	ets	2018 Q1	N(144, 1119)	144.
6	Business	ACT	<aggregated>	ets	2018 Q2	N(203, 2260)	203.
7	Business	ACT	<aggregated>	ets_adjusted	2018 Q1	N(157, 539)	157.
8	Business	ACT	<aggregated>	ets_adjusted	2018 Q2	N(214, 951)	214.
9	Business	NSW	Blue Mountains	ets	2018 Q1	N(20, 140)	19.7

Hierarchical and grouped time series

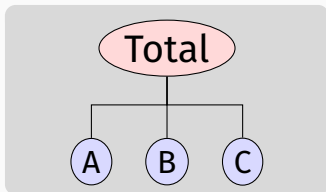
Every collection of time series with aggregation constraints can be written as

$$\mathbf{y}_t = \mathbf{S}\mathbf{b}_t$$

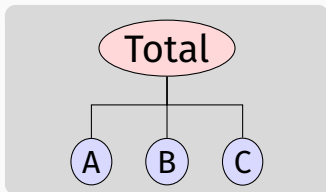
where

- \mathbf{y}_t is a vector of all series at time t
- \mathbf{b}_t is a vector of the most disaggregated series at time t
- \mathbf{S} is a “summing matrix” containing the aggregation constraints.

Hierarchical time series



Hierarchical time series

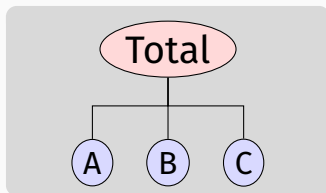


y_t : observed aggregate of all series at time t .

$y_{X,t}$: observation on series X at time t .

\mathbf{b}_t : vector of all series at bottom level in time t .

Hierarchical time series



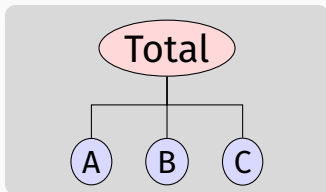
y_t : observed aggregate of all series at time t .

$y_{X,t}$: observation on series X at time t .

b_t : vector of all series at bottom level in time t .

$$\mathbf{y}_t = \begin{pmatrix} y_t \\ y_{A,t} \\ y_{B,t} \\ y_{C,t} \end{pmatrix} = \begin{pmatrix} 1 & 1 & 1 \\ 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} y_{A,t} \\ y_{B,t} \\ y_{C,t} \end{pmatrix}$$

Hierarchical time series



y_t : observed aggregate of all series at time t .

$y_{X,t}$: observation on series X at time t .

b_t : vector of all series at bottom level in time t .

$$\mathbf{y}_t = \begin{pmatrix} y_t \\ y_{A,t} \\ y_{B,t} \\ y_{C,t} \end{pmatrix} = \underbrace{\begin{pmatrix} 1 & 1 & 1 \\ 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix}}_{\mathbf{S}} \underbrace{\begin{pmatrix} y_{A,t} \\ y_{B,t} \\ y_{C,t} \end{pmatrix}}_{\mathbf{b}_t}$$

$$\mathbf{y}_t = \mathbf{S} \mathbf{b}_t$$

Forecasting notation

Let $\hat{\mathbf{y}}_n(h)$ be vector of initial h -step forecasts, made at time n , stacked in same order as \mathbf{y}_t .

Forecasting notation

Let $\hat{\mathbf{y}}_n(h)$ be vector of initial h -step forecasts, made at time n , stacked in same order as \mathbf{y}_t .
(In general, they will not “add up”.)

Forecasting notation

Let $\hat{\mathbf{y}}_n(h)$ be vector of initial h -step forecasts, made at time n , stacked in same order as \mathbf{y}_t .

(In general, they will not “add up”.)

Reconciled forecasts must be of the form:

$$\tilde{\mathbf{y}}_n(h) = \mathbf{S}\mathbf{G}\hat{\mathbf{y}}_n(h)$$

for some matrix \mathbf{G} .

Forecasting notation

Let $\hat{\mathbf{y}}_n(h)$ be vector of initial h -step forecasts, made at time n , stacked in same order as \mathbf{y}_t .
(In general, they will not “add up”.)

Reconciled forecasts must be of the form:

$$\tilde{\mathbf{y}}_n(h) = \mathbf{S}\mathbf{G}\hat{\mathbf{y}}_n(h)$$

for some matrix \mathbf{G} .

- \mathbf{G} extracts and combines base forecasts $\hat{\mathbf{y}}_n(h)$ to get bottom-level forecasts.
- \mathbf{S} adds them up

Optimal combination forecasts

Main result

The best (minimum sum of variances) unbiased forecasts are obtained when $\mathbf{G} = (\mathbf{S}'\Sigma_h^{-1}\mathbf{S})^{-1}\mathbf{S}'\Sigma_h^{-1}$, where Σ_h is the h -step base forecast error covariance matrix.

Optimal combination forecasts

Main result

The best (minimum sum of variances) unbiased forecasts are obtained when $\mathbf{G} = (\mathbf{S}'\Sigma_h^{-1}\mathbf{S})^{-1}\mathbf{S}'\Sigma_h^{-1}$, where Σ_h is the h -step base forecast error covariance matrix.

$$\tilde{\mathbf{y}}_n(h) = \mathbf{S}(\mathbf{S}'\Sigma_h^{-1}\mathbf{S})^{-1}\mathbf{S}'\Sigma_h^{-1}\hat{\mathbf{y}}_n(h)$$

Problem: Σ_h hard to estimate, especially for $h > 1$.

Solutions:

- Ignore Σ_h (OLS) [`min_trace(method='ols')`]
- Assume $\Sigma_h = k_h \Sigma_1$ is diagonal (WLS) [`min_trace(method='wls')`]
- Assume $\Sigma_h = k_h \Sigma_1$ and estimate it (GLS) [`min_trace(method='shrink')` (the default)]

Features

- Covariates can be included in initial forecasts.
- Adjustments can be made to initial forecasts at any level.
- Very simple and flexible method. Can work with *any* hierarchical or grouped time series.
- Conceptually easy to implement: regression of base forecasts on structure matrix.

Outline

1 Hierarchical and grouped time series

2 Forecast reconciliation

3 Example: Australian tourism

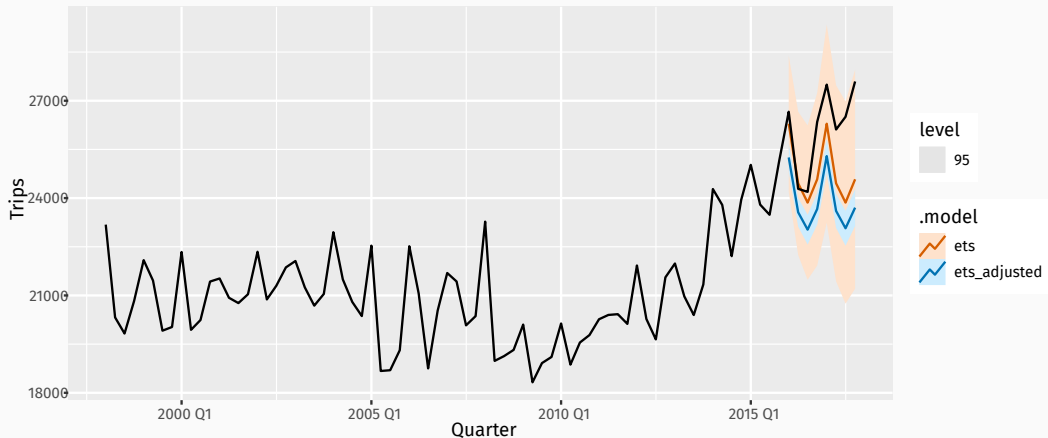
4 Lab Session 20

Example: Australian tourism

```
tourism_agg <- tourism |>
  aggregate_key(Purpose * (State / Region),
    Trips = sum(Trips)
  )
fc <- tourism_agg |>
  filter_index(. ~ "2015 Q4") |>
  model(ets = ETS(Trips)) |>
  reconcile(ets_adjusted = min_trace(ets)) |>
  forecast(h = "2 years")
```

Example: Australian tourism


```
fc |>  
  filter(is_aggregated(Purpose) & is_aggregated(State)) |>  
  autoplot(tourism_agg, level = 95)
```



Example: Australian tourism

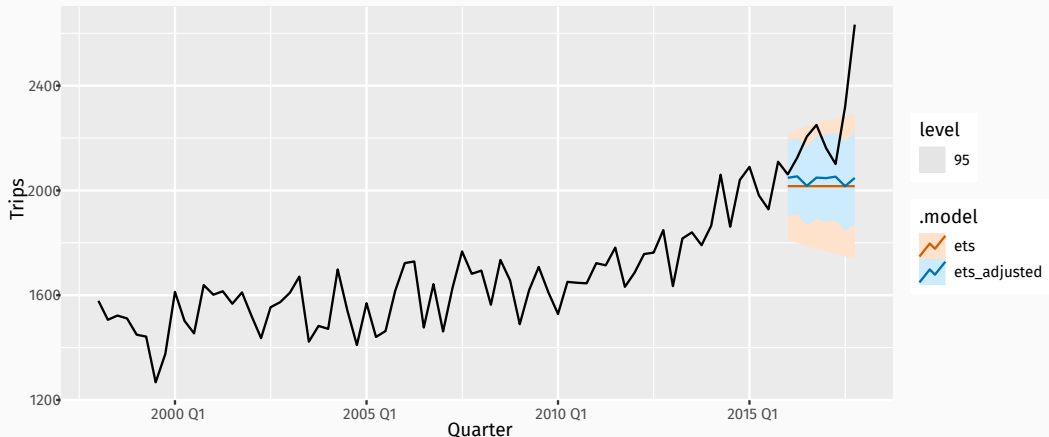
```
fc |>
```

```
  filter(is_aggregated(Purpose) & State == "VIC" & is_aggregated(Region)) |>  
  autoplot(tourism_agg, level = 95)
```




Example: Australian tourism

```
fc |>  
  filter(is_aggregated(Purpose) & Region == "Melbourne") |>  
  autoplot(tourism_agg, level = 95)
```



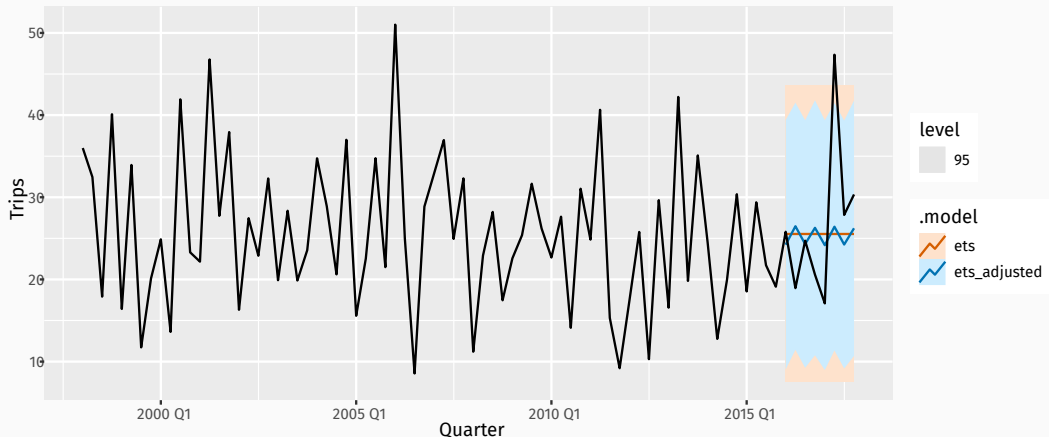
Example: Australian tourism

```
fc |>  
  filter(is_aggregated(Purpose) & Region == "Snowy Mountains") |>  
  autoplot(tourism_agg, level = 95)
```



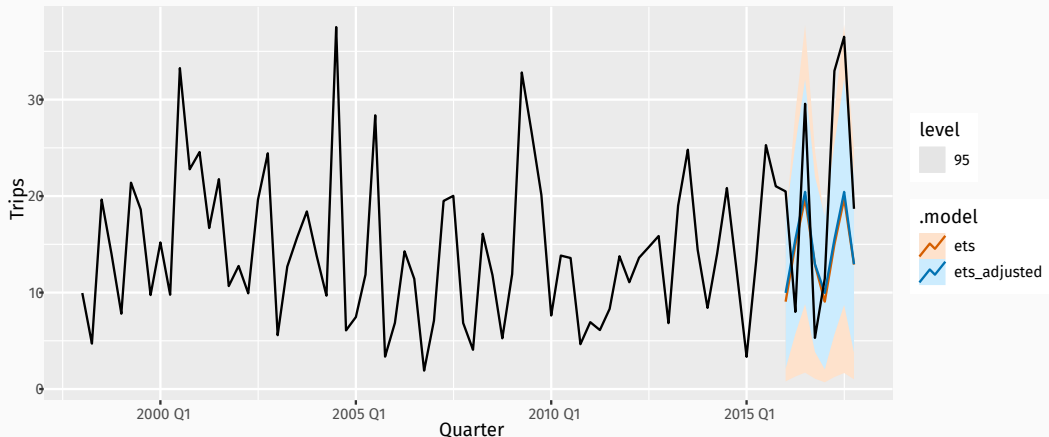
Example: Australian tourism

```
fc |>  
  filter(Purpose == "Holiday" & Region == "Barossa") |>  
  autoplot(tourism_agg, level = 95)
```



Example: Australian tourism

```
fc |>  
  filter(is_aggregated(Purpose) & Region == "MacDonnell") |>  
  autoplot(tourism_agg, level = 95)
```



Example: Australian tourism

```
fc <- tourism_agg |>
  filter_index(. ~ "2015 Q4") |>
  model(
    ets = ETS(Trips),
    arima = ARIMA(Trips)
  ) |>
  mutate(
    comb = (ets + arima) / 2
  ) |>
  reconcile(
    ets_adj = min_trace(ets),
    arima_adj = min_trace(arima),
    comb_adj = min_trace(comb)
  ) |>
  forecast(h = "2 years")
```

Forecast evaluation

```
fc |> accuracy(tourism_agg)
```

```
# A tibble: 2,550 x 13
```

	.model	Purpose	State	Region	.type	ME	RMSE	MAE	MPE	MAPE	MASE
	<chr>	<chr*>	<chr*>	<chr*>	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	arma	Business	ACT	Canberra	~ Test	35.9	45.7	35.9	16.9	16.9	0.938
2	arma	Business	ACT	<aggregated>	Test	35.9	45.7	35.9	16.9	16.9	0.938
3	arma	Business	NSW	Blue Mountains	~ Test	1.93	10.6	8.52	-18.0	48.6	0.644
4	arma	Business	NSW	Capital Country	~ Test	8.08	15.6	10.4	11.8	19.0	0.744
5	arma	Business	NSW	Central Coast	~ Test	10.0	14.5	10.8	26.9	32.2	0.982
6	arma	Business	NSW	Central NSW	~ Test	17.7	31.9	28.2	12.0	24.1	1.08
7	arma	Business	NSW	Hunter	~ Test	35.3	43.9	35.3	24.2	24.2	1.30
8	arma	Business	NSW	New England North W	~ Test	23.1	31.8	26.8	19.5	28.0	1.76
9	arma	Business	NSW	North Coast NSW	~ Test	24.8	40.1	36.8	11.5	28.5	1.38
10	arma	Business	NSW	Outback NSW	~ Test	6.87	11.0	7.76	13.7	16.5	0.571

```
# i 2,540 more rows
```

```
# i 2 more variables: RMSSE <dbl>, ACF1 <dbl>
```

Forecast evaluation

```
fc |>  
  accuracy(tourism_agg) |>  
  group_by(.model) |>  
  summarise(MASE = mean(MASE)) |>  
  arrange(MASE)
```

```
# A tibble: 6 x 2  
  .model      MASE  
  <chr>      <dbl>  
1 ets_adj    1.02  
2 comb_adj   1.02  
3 ets        1.04  
4 comb       1.04  
5 arima_adj  1.07  
6 arima      1.09
```

Outline

1 Hierarchical and grouped time series

2 Forecast reconciliation

3 Example: Australian tourism

4 Lab Session 20

Lab Session 20

- Prepare aggregations of the PBS data by Concession, Type, and ATC1.
- Use forecast reconciliation with the PBS data, using ETS, ARIMA and SNAIVE models, applied to all but the last 3 years of data.
- Which type of model works best?
- Does the reconciliation improve the forecast accuracy?
- Why doesn't the reconciliation make any difference to the SNAIVE forecasts?